

III-Logics

Introduction into Propositional Logics

Logik für Erklärbare KI: Technische Einführung in das ENEXA Projekt

Maria Schnödt-Fuchs, Alex Goëßmann

Funded by the
European Union



15.+16. July, 2024

Let there be a system described by three binary variables (being **True** or **False**):

- ▶ **Wet**: Observation, that the street is wet.
- ▶ **Rained**: It had rained.
- ▶ **Sprinkler**: The sprinkler had been on.

Example of Reasoning

When we know that the formulas

- ▶ **Rained** \Rightarrow **Wet** : Whenever it had rained, the street is wet.
- ▶ **Rained**: It had rained.

hold (i.e. = **True**), then it follows that the street is wet (i.e. **Wet** = **True**).

Logic is build upon **Syntax**

- ▶ How to formulate logical formulas?
- ▶ Example: The formula **Rained** \Rightarrow **Wet** represents the implication, that the street is wet whenever it had rained.
- ▶ Propositional Logics: The variables (or atomic formulas) are combined with logical connectives $\{\neg, \wedge, \vee, \Rightarrow, \Leftrightarrow\}$ to build generic formulas.

and **Semantics**

- ▶ What do these formulas describe?
- ▶ Example: From the formula **Rained** \wedge (**Rained** \Rightarrow **Wet**) the formula **Wet** follows.
- ▶ Propositional Logics: Each formula describes possible worlds, which are called models of a formula. Comparing the models defines the **entailment relation**.

Model-theoretic semantics: Matrices represent formulas with two variables

We mark the models (the possible worlds) of a formula, and choose an order by

- ▶ **Wet** is False on the first column and True on the second
- ▶ **Rained** is False on the first row and True on the second

Then the model-theoretic semantics of $\text{Rained} \Rightarrow \text{Wet}$ is represented by the matrix:

$$(\text{Rained} \Rightarrow \text{Wet}) = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix} \quad (1)$$

Entailment by comparing the position of ones

We summarize

$$(\text{Rained} \Rightarrow \text{Wet}) = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix} \quad \text{and} \quad \text{Rained} = \begin{bmatrix} 1 & 0 \\ 1 & 0 \end{bmatrix} \quad (2)$$

into

$$\text{Rained} \wedge (\text{Rained} \Rightarrow \text{Wet}) = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}. \quad (3)$$

Comparing with

$$\text{Wet} = \begin{bmatrix} 1 & 1 \\ 0 & 0 \end{bmatrix}. \quad (4)$$

we notice, that every model of $\text{Rained} \wedge (\text{Rained} \Rightarrow \text{Wet})$ is a model of Wet . We say that Wet is **entailed** and denote

$$\left(\text{Rained} \wedge (\text{Rained} \Rightarrow \text{Wet}) \right) \models \text{Wet}. \quad (5)$$

Model-theoretic semantics: Tensors represent multiple variables

Tensors are a way to order all possible worlds in a system with multiple variables:

- ▶ Each variable is associated with an **axis**.
- ▶ For each assignment to a variable we find a **coordinate**

Example of three variables:

We associate the variable **Sprinkler** with a third direction

$$(\text{Rained} \Rightarrow \text{Wet}) \wedge \text{Sprinkler} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}$$

Factored System

A factored system is a set of d variables X_k indexed by $k \in [d]$ each having 2 assignments, which interact with each other.

The total number of states of a factored system is

$$2^d.$$

Curse of dimensionality

When reasoning about a factored system with many variables (i.e. d is large), the naive enumeration of all models to formulas is infeasible.

The curse of dimensionality can be mitigated by smart distributed representations of the model marking tensors.


Strategy: Store the semantics of formulas f_1 , f_2 and $f_1 \circ f_2$ at each decomposition of the formula.

Distributed Representation

The semantics of complex formulas is then stored by a set of semantics of the used logical connectives $\circ \in \{\neg, \wedge, \vee, \Rightarrow, \Leftrightarrow\}$.

A **graphical notation** depicts generic tensors:

- ▶ Tensors are rectangular boxes
- ▶ Each axis of a tensor is drawn by a leg to the box




Encoding Logics using Tensors

Boolean states are represented by one-hot encodings

$$e_{\text{True}} = e_1 = \begin{bmatrix} 0 \\ 1 \end{bmatrix} \quad \text{and} \quad e_{\text{False}} = e_0 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}.$$

The truth tables of logical connectives $\circ \in \{\wedge, \vee, \Rightarrow\}$ are represented by tensors $\theta^\circ \in \mathbb{R}^{2 \times 2 \times 2}$ defined as

$$\theta^\circ = \sum_{f_1, f_2 \in \{\text{True}, \text{False}\}} e_{f_1} \otimes e_{f_2} \otimes e_{f_1 \circ f_2}$$



Example: Tensor Core to logical disjunction \vee has the coordinates

$$\theta_{1,:,:}^\vee = \begin{bmatrix} 0 & 1 \\ 1 & 1 \end{bmatrix} \quad \text{and} \quad \theta_{0,:,:}^\vee = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}.$$

Encoding of Generic Formulas

The semantics of a formula f is a map


$$f : \bigtimes_{k \in [d]} [2] \rightarrow \{0, 1\}.$$

Given a formula f we call a world index by $(i_0, \dots, i_{d-1}) \in \bigtimes_{k \in [d]} [2]$ a model of f , if $f(i_0, \dots, i_{d-1}) = 1$.

formulas are encoded by tensors

$$\theta^f = \sum_{i_0, \dots, i_{d-1} \in [2]} \left(\bigotimes_{k \in [d]} e_{i_k} \right) \otimes e_{f(i_0, \dots, i_{d-1})}$$

depicted as




Representation of Complex Formulas


The semantics of complex formulas are retrieved by **contractions** of their connective semantics, which are summations of tensor coordinates among shared axes.

- ▶ Choose distributed representation to avoid contractions
- ▶ Only execute those contractions required by reasoning

Contractions can be depicted graphically by a **Tensor Network**:



b)



The semantics of logics enable the calculation of logical consequences, framed logical inference.

Definition (Entailment)

A propositional formula f is entailed by a formula \mathcal{KB} , if for each state i_0, \dots, i_{d-1} with $\mathcal{KB}(i_0, \dots, i_{d-1}) = 1$ we have $f(i_0, \dots, i_{d-1}) = 1$. We denote in this case

$$\mathcal{KB} \models f .$$

Approaches to automate the decision of entailment:


- ▶ Proof-theoretic approaches: Inference rules such as Modus ponens
- ▶ Model-theoretic approaches: Check whether $\mathcal{KB} \wedge \neg f$ has a model.

Theorem (Contraction Criterion for Entailment)

We have $\mathcal{KB} \models f$ (respectively $\mathcal{KB} \models \neg f$) if and only if

$$\mathcal{C} \left(\{\mathcal{KB}, \theta^f\}, \{X_f\} \right) \parallel e_1 \quad (\text{respectively } \mathcal{C} \left(\{\mathcal{KB}, \theta^f\}, \{X_f\} \right) \parallel e_0).$$

We depict this condition by



Model-theoretic Semantics by Tensor Networks


- ▶ **Tensors encode the models** of propositional formulas by their nonzero coordinates
- ▶ **Networks eliminate redundancies** when representing similar formulas

This is an implementation of the **neural paradigm**, resulting in

- ▶ Effective representation: Latent variables by logical formulas
- ▶ Differential parametrization: Multilinearity with respect to each core tensor

Implementations in tnreason : Subpackage encoding

The subpackage encoding is dedicated to the implementation of logics.



Propositional formulas have three representations specified in the subpackage encoding :

- ▶ **Syntax** of formulas f is stored in a script language $S(f)$ based on nested lists of strings
- ▶ **Semantics** of formulas is stored by tensor networks of connective cores
- ▶ Random variable representing the formula satisfaction

In tntreason , propositional syntax is represented by [nested lists of strings](#), for example:

```
["and", ["not", "Rained"], ["imp", "Rained", "Wet"]]
```

Connective have a defined representation as strings:

Unary connective \circ	$S(\circ)$
\neg	"not"
()	"id"

Binary connective \circ	$S(\circ)$
\wedge	"and"
\vee	"or"
\Rightarrow	"imp"
\oplus	"xor"
\Leftrightarrow	"eq"

Any other string is interpreted by an atomic formula (a so-called propositional symbol).

Having strings representing

- ▶ Connectives (by defined representations)
- ▶ Atomic Formulas (all other strings)

we represent formulas $f_1 \circ, f_2$

$[S(\circ), S(f_1), S(f_2)]$

where we apply the conventions

- ▶ Connectives are at the 0th position in each list
- ▶ Further entries are either atoms as strings or encoded formulas itself

Examples of representations

Atomic variable Rained by

$S(\text{Rained}) = \text{"Rained"}$

Negative literal $\neg\text{Rained}$ by

$[\text{"not"}, \text{"Rained"}]$

Horn clause $(\text{Rained} \Rightarrow \text{Wet})$ by

$[\text{"imp"}, \text{"Rained"}, \text{"Wet"}]$

Knowledge Base $(\neg\text{Rained}) \wedge (\text{Rained} \Rightarrow \text{Wet})$ by

$[\text{"and"}, [\text{"not"}, \text{"Rained"}], [\text{"imp"}, \text{"Rained"}, \text{"Wet"}]]$

The recursive structure of the nested lists $S(f)$ is exploited in finding tensor network representations of f . The function

```
encoding.create_raw_cores()
```

creates the connective cores to $S(f)$ by

- ▶ When $S(f)$ a list, create $\theta^{S(f)[0]}$ and add to the tensor networks created by recursive calls to the subformulas $S(f)[1 :]$
- ▶ Return empty list, when $S(f)$ is a string

Then we have

$$f = \mathcal{C} (\{encoding.create_raw_cores(S(f))\} \cup \{e_1\}, \{X_0, \dots, X_{d-1}\})$$