

Tensor networks for neuro-symbolic AI

Alex Goessmann

joint work with Janina Schütte, Maximilian Fröhlich and Martin Eigel

arXiv : 2601.15442

Weierstraß Institute for Applied Analysis and Stochastics (WIAS)

NESYRE Workshop
September 15, 2026

- ▶ Tensor networks for neuro-symbolic AI
- ▶ Maximum entropy principles
- ▶ Summary and Outlook

- ▶ Tensor networks for neuro-symbolic AI
- ▶ Maximum entropy principles
- ▶ Summary and Outlook


Toy accounting rule system

Assume you want to book one of two accounts A1 and A2 depending on a feature F of an invoice, based on two rules:


- ▶ Exactly one account must be booked.
- ▶ If feature F is present on the invoice, the account A1 is typically booked.
- ▶ Accounting rules are formalised in propositional logics, triggered by receipts properties.
- ▶ These rules can either be hard restrictions, or prone to exceptions.

We apply tensor networks for neuro-symbolic use cases like this!

The variables and main relations of the DATEV accounting ontology¹:



Coarse Visualization



Detailed Visualization

¹<https://github.com/EnexaProject/enexa-use-case-1-business-software-solutions>

Tensors: Natural description of factored representations


Factored representations are collection of variables $X_{[d]}$, each with a finite set of states $[m_k] = \{0, \dots, m_k - 1\}$ [?]. The set of states is then the cartesian product

$$\times_{k \in [d]} [m_k]$$

Knowledge about the system, for example the probability of each state, is a **tensor**

$$\tau [X_{[d]}] : \times_{k \in [d]} [m_k] \rightarrow \mathbb{R}$$

We represent them by diagrams with legs decorated by the variables:



Necessity

$$\tau^{\mathcal{G}} [X_{\mathcal{V}} = x_{\mathcal{V}}] = \prod_{e \in \mathcal{E}} \tau^e [X_e = x_e]$$

Typically: Applied in a Black Box manner Here: We relate decompositions to the logical and probabilistic structure


Neural sparsity

Functions of interest decompose into simple functions (neurons).

Capture by tensor networks: **Basis encodings**

$$\beta^f [Y_{[p]}, X_{[d]}] = \sum_{X_{[d]} \in \times_{k \in [d]} [m_k]} \epsilon_{f(X_{[d]})} [Y_{[p]}] \otimes \epsilon_{X_{[d]}} [X_{[d]}]$$

Example: The basis encoding of $f [X_{[3]}] = (X_0 \vee X_1) \wedge \neg X_2$ is decomposed as:




Logical syntax translates in tensor network formats

Graphical models: Directly tensor networks

Tensor network graphs encode conditional independences

Computation-Activation Networks (CompActNets)

A novel tensor network architecture captures logical and probabilistic models (see [?]):



Activation of a statistic by an activation core, for **knowledge bases**

$$\alpha [Y_{[p]}] = \bigotimes_{\ell \in [p]} \begin{pmatrix} 0 \\ 1 \end{pmatrix} [Y_{\ell}]$$

and for **exponential families**

$$\alpha [Y_{[p]}] = \bigotimes_{\ell \in [p]} \begin{pmatrix} 1 \\ \exp [\theta [L = \ell]] \end{pmatrix} [Y_{\ell}]$$


Computation of a statistic

$$t : \times_{k \in [d]} [m_k] \rightarrow \times_{\ell \in [p]} [p_{\ell}]$$

by basis encoding

Toy accounting system as a CompActNet

We represent the toy accounting system as a Computation-Activation Network:




- ▶ The hard activation vector $\begin{pmatrix} 0 \\ 1 \end{pmatrix} [y_{A1 \oplus A2}]$ ensures that the first rule is always true.
- ▶ The soft activation vector $\begin{pmatrix} 1 \\ \exp[\theta] \end{pmatrix} [y_{F \Rightarrow A1}]$ introduces the trainable parameter $\theta \in \mathbb{R}$ controlling the satisfaction probability of the second rule.

Main inference principle


Model inference amounts to tensor network contractions.

- Probabilistic inference: Contractions are marginal probabilities.
- Logical inference: Contractions are model counts, which decide entailment.

Example: Distribution of X_{A1} and X_{A2} conditioned on $X_F = 0$ is parallel to



Tensor networks provide a unifying language to capture the different paradigms of AI [?]:



- ▶ Tensor networks for neuro-symbolic AI
- ▶ Maximum entropy principles
- ▶ Summary and Outlook

Shannon Entropy: Natural measure of structure (negative information content) of a distribution.

Maximum entropy principle

Among those distributions reproducing certain characteristics of observed data, choose the one with the least structure (=maximum entropy).

Formally the maximum entropy problem is

$$\operatorname{argmax}_{\mathbb{P}[X_{[d]}]} \mathbb{H}[\mathbb{P}[X_{[d]}]] \quad \text{s.t.} \quad \langle \mathbb{P}[X_{[d]}], \sigma^t[X_{[d]}, L] \rangle_{[L]} = \mu[L] .$$

Exponential families are maximum entropy distributions

Not all maximum entropy distributions are in the exponential family.
 Characterization by the face of the mean polytope to which the mean parameter belongs.

Realizable mean parameters form a convex polytope, the **mean polytope** of the statistic t :


Exponential families can only realize the relative interior of the mean polytope.

What are the solutions to the mean parameters on the boundary of the mean polytope?

Consider the statistic


$$f_0[X_0, X_1] = (X_0 \wedge X_1) \quad , \quad f_1[X_0, X_1] = (X_0 \Rightarrow X_1)$$

The mean polytope of this statistic is



We notice: The blue face cannot be represented by an elementary tensor.

Representation of the maximum entropy distribution for a mean parameter in the blue face:



- ▶ Tensor networks for neuro-symbolic AI
- ▶ Maximum entropy principles
- ▶ Summary and Outlook

Tensor networks

- ▶ Tensors appear naturally as factored system representations
- ▶ **Sparsity principle**: Decomposition into tensor networks
- ▶ **Inference principle**: Tensor network contractions

Implementations in python packages:

- ▶ `tnreason`²: Tensor network based reasoning
- ▶ `qcreason`³: Quantum inference for contractions

²<https://pypi.org/project/tnreason/>

³<https://pypi.org/project/qcreason/>

Thank you very much for your attention!